

# Class 7: Machine Learning I

Youn Soo Na (PID: A1704731)

Today we are going to learn how to apply different machine learning methods, beginning with clustering:

The goal here is to find groups/clusters in your input data.


First, I will make up some data with clear groups. For this I will use the `rnorm()` function:

```
rnorm(10)
```


```
[1] -1.11017313  0.32945160 -0.12018342 -0.32179992  1.77466785  0.32087312  
[7]  0.38968077  0.62894053  0.01826351  0.07473966
```

```
hist( rnorm(10000, mean = 3) )
```

**Histogram of `rnorm(10000, mean = 3)`**




```
n <- 10000
x <- c(rnorm(n,-3), rnorm(n, +3))
hist(x)
```



```
n <- 30
x <- c(rnorm(n,-3), rnorm(n, +3))
y <- rev(x)
z <- cbind(x,y)
head(z)
```

	x	y
[1,]	-4.445269	2.840753
[2,]	-3.521256	2.543306
[3,]	-3.267614	3.564238
[4,]	-3.054634	3.114294
[5,]	-4.411724	2.683313
[6,]	-3.557461	2.177102

```
plot(z)
```



Use the `kmeans()` function setting `k` to 2 and `nstart=20`

Inspect/print the results

Q. How many points are in each cluster?

Q. What ‘component’ of your result object details - cluster size? - cluster assignment/membership? - cluster center?

```
km <- kmeans(z, centers = 2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	2.942853	-2.943101
2	-2.943101	2.942853

Clustering vector:

[illegible]


```
[1] 69.3442 69.3442
(between_SS / total_SS = 88.2 %)
```

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

	x	y
1	2.942853	-2.943101
2	-2.943101	2.942853


Q. Plot x colored by the kmeans cluster assignment and add cluster centers as blue points.

```
plot(z, col = c("red", "blue"))
```




R will re-cycle the shorter color vector to be the same length as the longer (number of data points) in z.

```
plot(z, col = km$cluster)
```



We can use the `points()` function to add new points to an existing plot. . .


```
plot(z, col = km$cluster)
points(km$centers, col="blue", pch=33, cex = 3)
```



```
# max pch is 127
```

Q. Can you run kmeans and ask for 4 clusters please and plot the results like we have done above?

```
km4 <- kmeans(z, centers = 4)
plot(z, col = km4$cluster)
points(km4$centers, col="blue", pch=33, cex = 3)
```



## Hierarchical Clustering

Let's take our same made-up data z and see how hclust works.

First we need a distance matrix of our data to be clustered.


```
d <- dist(z)
hc <- hclust(d)
hc
```

Call:

```
hclust(d = d)
```

```
Cluster method   : complete  
Distance         : euclidean  
Number of objects: 60
```

```
plot(hc)
```



I can get my cluster membership vector by “cutting the tree” with the `cutree()` function like so


```
grps <- cutree(hc, h=8)  
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 1 1 1 1 1 1 3 3 3 3 3 3 2 3  
[39] 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
```

Can you plot `z` colored by out `hclust` results?

```
plot(z, col=grps)
```






## PCA of UK Food Data

Read data from the UK on food consumption in different parts of the UK

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
head(x)
```


	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



A so-called “Pairs” plot can be useful for small datasets like this

```
pairs(x, col=rainbow(10), pch=16)
```



It's hard to see structure and trends in even this small data-set. How will we ever do this when we have big data-sets with 1,000s or 10s of thousands of things we are measuring. . .

## PCA to the rescue

Let's see how PCA deals with this dataset. So main function in base R to do PCA is called `prcomp()`

```
#transpose t()
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Let's see what is inside this `pca` object that we created from running `prcomp()`

```
attributes(pca)
```


```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
[1] "prcomp"
```

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
plot(pca$x[,1], pca$x[,2], col = c("black", "red", "blue", "darkgreen"), pch=16,
     xlab="PC1 (67.4%)", ylab = "PC2 (29.0%)")
text(pca$x[,1], pca$x[,2], labels=c("England", "Wales", "Scotland", "N.Ireland"), col=c("black", "red", "blue", "darkgreen"))
```



### Variable Loadings Plot

```
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```

